

# **WirelessLab WS 2016/17**

## **Homework 6: Medium Utilization**

**Group 6**

**Gasper Kojek, Jens Klein**

# Question 1

## a) Data throughput with and without RTS/CTS

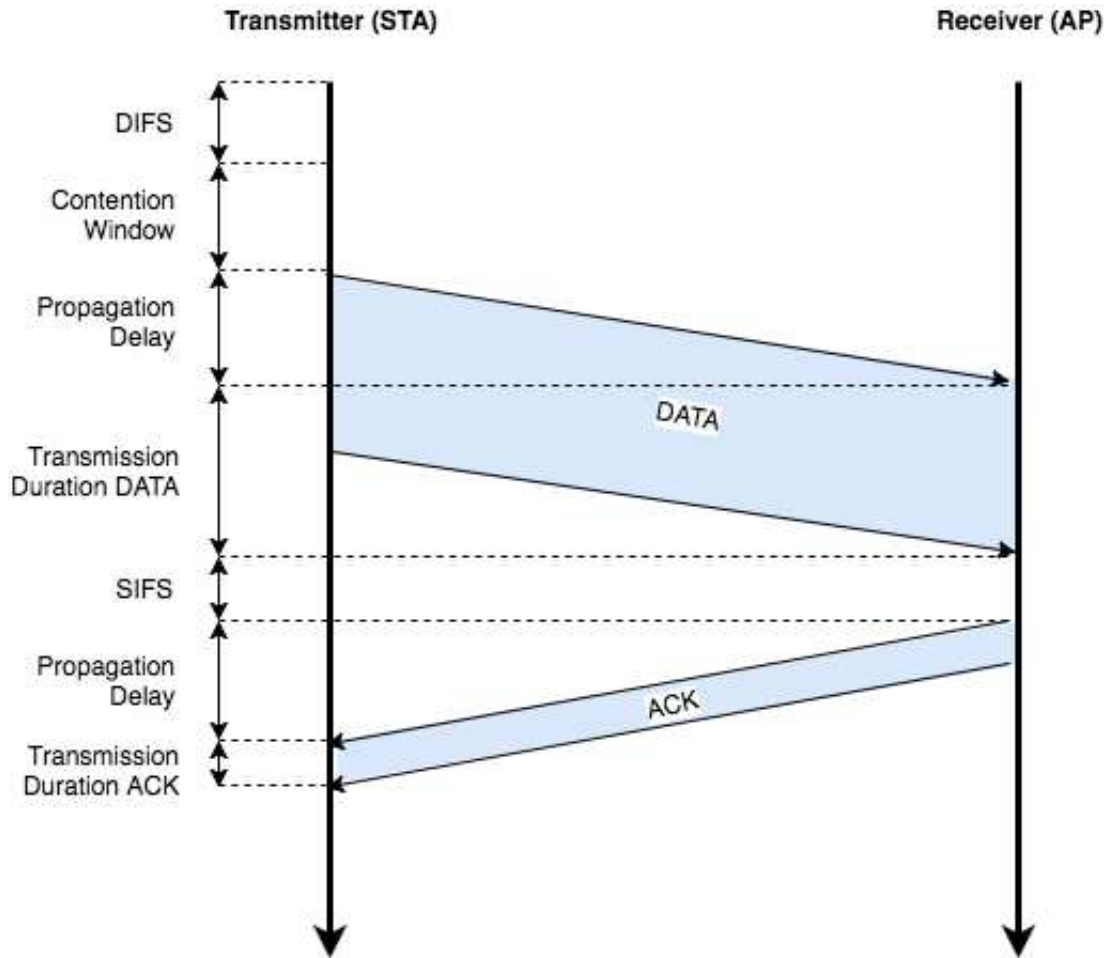
---

The following figure gives an overview of the data flow from a STA to an AP without using RTS/CLS mechanism. A few general assumptions were made:

- no loss or collision during transmission
- no interference
- no hidden or exposed terminals
- no fragmentation

### Without RTS/CTS

Flow diagram:



Defintions:

- DIFS time:  $t_{DIFS} = 34 \mu s$
- Slot Time:  $t_{ST} = 9 \mu s$
- Maximum backoff slots:  $b_{max} = 15$
- Random backoff:  $RB = \{n : n \text{ is an integer; and } 1 \leq n \leq b_{max}\}$
- Expected backoff:  $b_{expt} = \frac{b_{max}+1}{2} = 8$
- Contention window:  $t_{CW} = b_{expt} \cdot t_{ST} = 72 \mu s$
- Propagation delay:  $t_{pd} = 1 \mu s$
- SIFS time:  $t_{SIFS} = 16 \mu s$
- PHY layer overhead =  $t_{phy} = 20 \mu s$
- OFDM symbol duration =  $t_{ODFM} = 4 \mu s$
- MAC layer data payload =  $d_{mpay} = 1452 B$
- MAC header size =  $d_{mhead} = 28 B$
- MAC ack size =  $d_{mack} = 14 B$
- PHY Layer transmission rate:  $r = 54 \text{ Mbps} = 7.077888 \frac{B}{\mu s}$
- Transmission duration data:  $t_{data} = t_{phy} + t_{ODFM} + \frac{d_{mhead} + d_{mpay}}{r} \approx 243 \mu s$

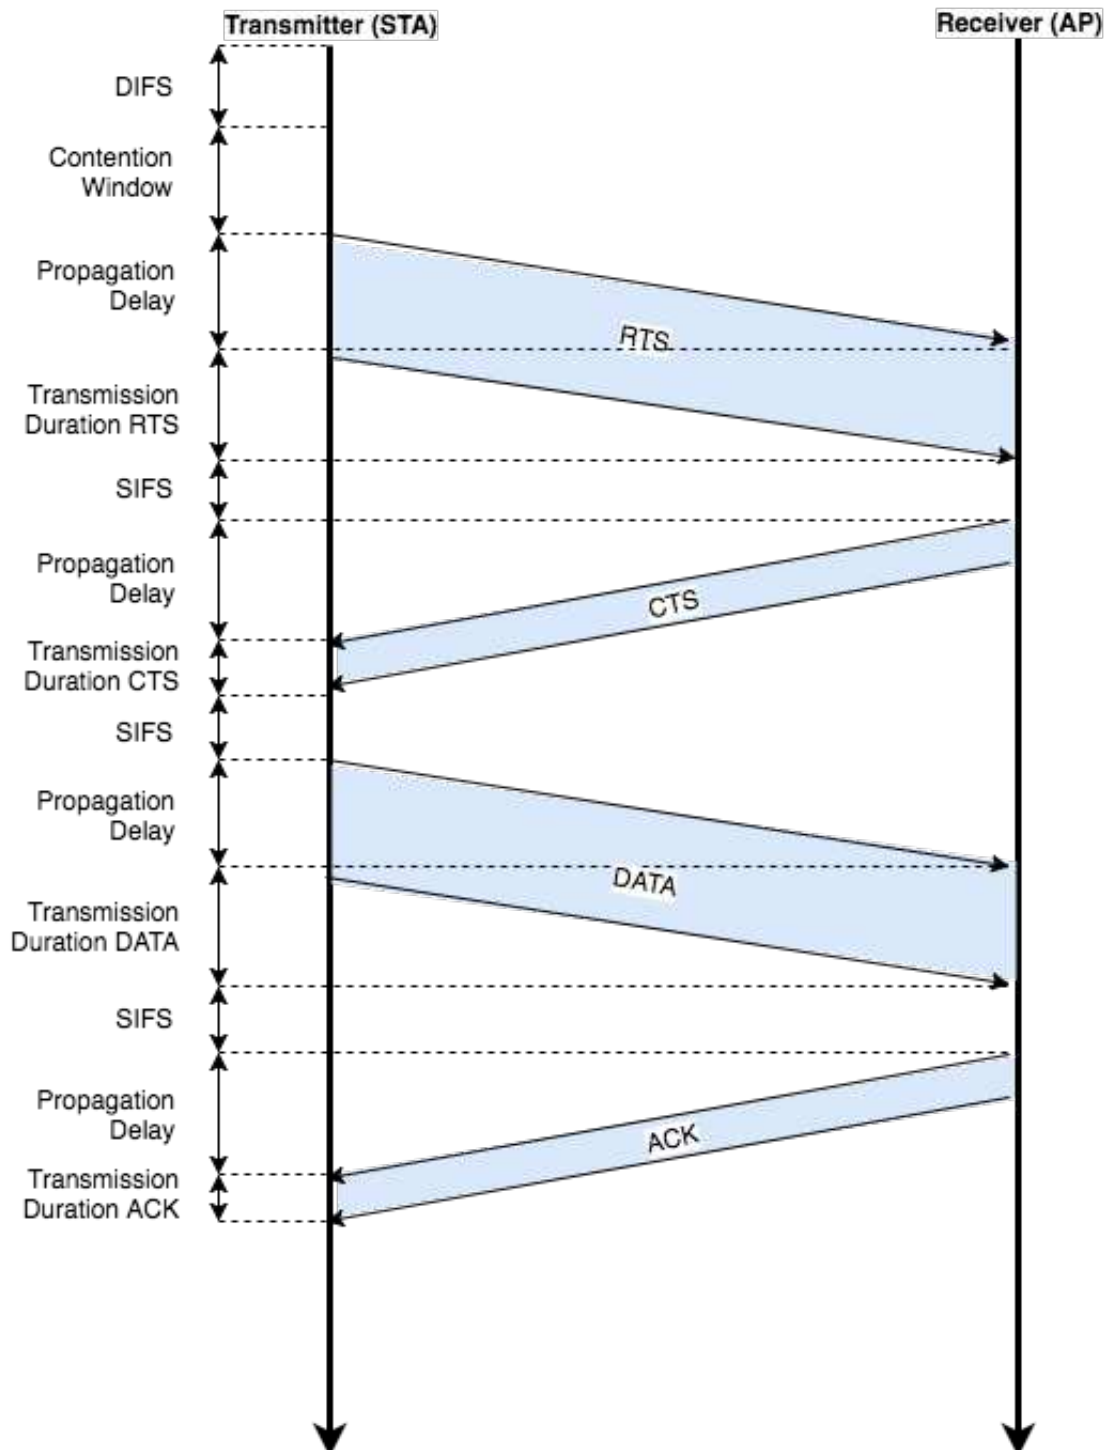
- Transmission duration ack:  $t_{ack} = t_{phy} + t_{ODFM} + \frac{d_{mhead} + d_{mack}}{r} \approx 30 \mu s$
- Total Time:  $t_{total} = t_{DIFS} + t_{CW} + 2 \cdot t_{pd} + t_{data} + t_{SIFS} + t_{ack} \approx 397 \mu s$

Calculated Transmission Rate:

$$r_{act} = \frac{d_{mpay}}{t_{total}} \approx \frac{1452 B}{397 \mu s} \approx \frac{1.11 \cdot 10^{-2} Mbit}{3.97 \cdot 10^{-4} s} \approx 28.0 Mbps$$

## With RTS/CTS

Flow diagram:



Additional Definitions to the previous:

- CTS size:  $d_{cts} = 14 B$
- RTS size:  $d_{rts} = 20 B$
- Transmission duration CTS:  $t_{cts} = t_{phy} + t_{ODFM} + \frac{d_{mhead} + d_{cts}}{r} \approx 30 \mu s$
- Transmission duration CTS:  $t_{cts} = t_{phy} + t_{ODFM} + \frac{d_{mhead} + d_{rts}}{r} \approx 31 \mu s$
- Total Time:  $t_{total} = t_{DIFS} + t_{CW} + 4 \cdot t_{pd} + t_{rts} + t_{cts} + t_{data} + 3 \cdot t_{SIFS} + t_{ack} \approx 493 \mu s$

Calculated Transmission Rate:

$$r_{act} = \frac{d_{mpay}}{t_{total}} \approx \frac{1452 B}{493 \mu s} \approx \frac{1.11 \cdot 10^{-2} Mbit}{4.93 \cdot 10^{-4} s} \approx 22.5 Mbps$$

## Conclusion

With RTS/CTS disabled there is a theoretical transmission rate of about 28.0 Mbps. With RTS/CTS enabled it is about 22.5 Mbps. That is 20 % less throughput. The reason are the additional frames for the handshake and the additional propagation delays and SIFS.

## b) Data throughput with and without RTS/CTS

---

Terminology:

- N6 = Node 6
- N15 = Node 15
- ST = Stepping Stone

## Setup

- N6 is set as the AP?

N6: `iw wlan0 info | grep type`

Output: `type AP`

- Get N6 IP address N6: `ifconfig wlan0 | grep "inet addr"`

Output: `inet addr:172.17.5.10 Bcast:172.17.5.255 Mask:255.255.255.0`

- N15 is set as client and connected to AP of N6

N15: `ping -I wlan0 172.17.5.10`

Output (trunc): `64 bytes from 172.17.5.10: seq=0 ttl=64 time=0.898 ms`

- Enable RTS/CTS on N15:

N15: `iw phy phy0 set rts 100`

- Set bitrates on both interfaces:

N15: `iw wlan0 set bitrates legacy-2.4 54.0` N6:

`iw wlan0 set bitrates legacy-2.4 54.0`

- Set tx power on client

N15: `iw wlan0 set txpower fixed 30.0`

- Review settings on client

N15: `iwinfo`

Output:

```
wlan0 ESSID: "group06_ap"
      Access Point: 00:1B:B1:07:DB:9B
      Mode: Client  Channel: 11 (2.462 GHz)
      Tx-Power: 30 dBm  Link Quality: 70/70
      Signal: -38 dBm  Noise: -96 dBm
      Bit Rate: 54.0 MBit/s
      Encryption: none
      Type: nl80211  HW Mode(s): 802.11abg
      Hardware: 168C:0013 185F:1012 [Generic MAC80211]
      TX power offset: unknown
      Frequency offset: unknown
      Supports VAPs: yes  PHY name: phy0
```

- Start iperf server

N6: `iperf -s -u`

- 
- Start client with 1400 B UDP datagrams CTS/RTS Threshold 100 B

```
for i in `seq 10`; do
    iperf -c 172.17.5.10 -u -b 54M -t 30 -l 1400
    sleep 2
done
```

- Start client with 200 B UDP datagrams CTS/RTS Threshold 100 B

```
for i in `seq 10`; do
    iperf -c 172.17.5.10 -u -b 54M -t 30 -l 200
    sleep 2
done
```

- disable CTS/RTS

N15: `iw phy phy0 set rts off`


- Start client with 1400 B UDP datagrams CTS/RTS off

```
for i in `seq 10`; do
    iperf -c 172.17.5.10 -u -b 54M -t 30 -l 1400
    sleep 2
done
```

- Start client with 200 B UDP datagrams CTS/RTS off

```
for i in `seq 10`; do
    iperf -c 172.17.5.10 -u -b 54M -t 30 -l 200
    sleep 2
done
```

## Boxplots



## Conclusion

The figure above shows boxplots from four experiments which are the combinations of CTS/RTS off and on, and UDP Packet Size of 1400 Bytes and 200 Bytes. The UDP data sent is fixed on all experiments (54

Mbps). We expected RTS/CTS enabled to have a smaller throughput with same packet lengths because it introduces additional MAC frames that do not transport actual payload but control the traffic. Also, in theory, RTS/CTS enabled results in a more stable throughput, because it reduced the risk of frame loss due to collisions.

RTS/CTS disabled produce the best throughput respectively on the same UDP size but are also more spread visible at the CIs. Exactly as expected.

What is different from our expectations is the difference between the calculated theoretical and the actual throughput. For UDP with about 1400 B size we calculated 28 Mbps. But we measured around 18 Mbps (-36% less) for CTS/RTS disabled. With this feature enabled we measured around 8 Mbps (-65% than calculated). In general the actual throughput is less because of other traffic on the medium like beacons and other control and management frames as well as interference.

A possible reason for CTS/RTS-enabled-throughput is even slower (compare percentages in previous paragraph) is that there might be exposed terminals that prevents the client from sending.



# Question 2

## a) Setup

---

- Reset bitrates on client

N15: `iw wlan0 set bitrates`

- Turn RTS/CTS off on client

N15: `iw phy phy0 set rts off`

- Prepare a script to run on node 6: `capture_manager.sh`

```

# !/usr/bin ash

STARTTIME=`date +%s`
DUMPFIL="tcpdump_{$STARTTIME}.cap"
SURVEYFILE="surveydump_{$STARTTIME}.dump"

echo ""
echo "Starting capture manager"
echo "Make sure you started nc on SteppingStone with command:  nc -l -p 8080 >
  filename.cap"

# reset survey counter
echo "  Resetting counter of survey dump"
echo "1" > "/sys/kernel/debug/ieee80211/phy0/ath5k/reset"

echo "  Starting tcpdump in the background ..."
# tcpdump -i wlan1 -w- > $DUMPFIL &
tcpdump -i wlan1 -w- | nc 172.17.3.1 8080 &
TCPDUMP_PID=`ps | grep "tcpdump -i wlan1" | grep -v grep | awk '{print $1}'`
echo "  tcpdump started with PID $TCPDUMP_PID"

for i in `seq 50`; do
  date +%s >> $SURVEYFILE
  iw dev wlan0 survey dump >> $SURVEYFILE
  sleep 0.1
done

echo "Starting iperf server"
iperf -s -u > /dev/null 2>&1 &
IPERF_PID=`ps | grep "iperf -s -u" | grep -v grep | awk '{print $1}'`
echo "iperf started with PID $IPERF_PID"

for i in `seq 200`; do
  date +%s >> $SURVEYFILE
  iw dev wlan0 survey dump >> $SURVEYFILE
  sleep 0.1
done

kill $IPERF_PID
sleep 1
kill $TCPDUMP_PID

```

- Commands for the test:

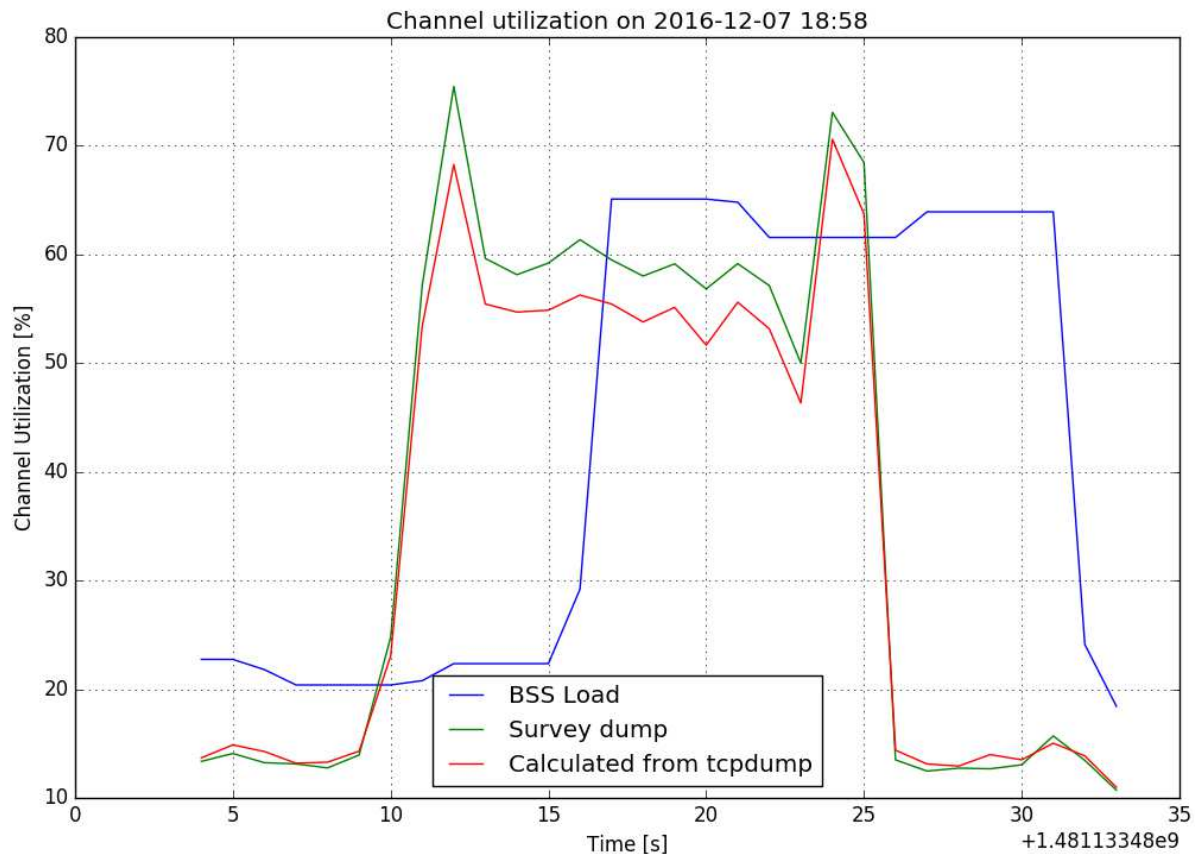
ST: `nc -l -p 8080 > filename.cap`

N6: `./capture_manager.sh`

N15: (after we see that iperf has started on N6): `iperf -c 172.17.5.10 -u -b 54M -t 15`

## c) Medium Utilization over time

### Plot



### Conclusion

We can see a slight difference between values received from `iw wlan0 survey dump` and values calculated manually from tcpdump trace using data rate and length of frames, which can be attributed to the fact that in we omitted times for SIFS, DIFS, contention window and propagation delay. We omitted those times because, if we wanted to include them in our calculations the right way, we would have to check for all the different types of frames when which length applies. Doing it this way also lets us know what percentage of medium it actually takes to transmit frames. We did however include the ACK time for UDP packets.

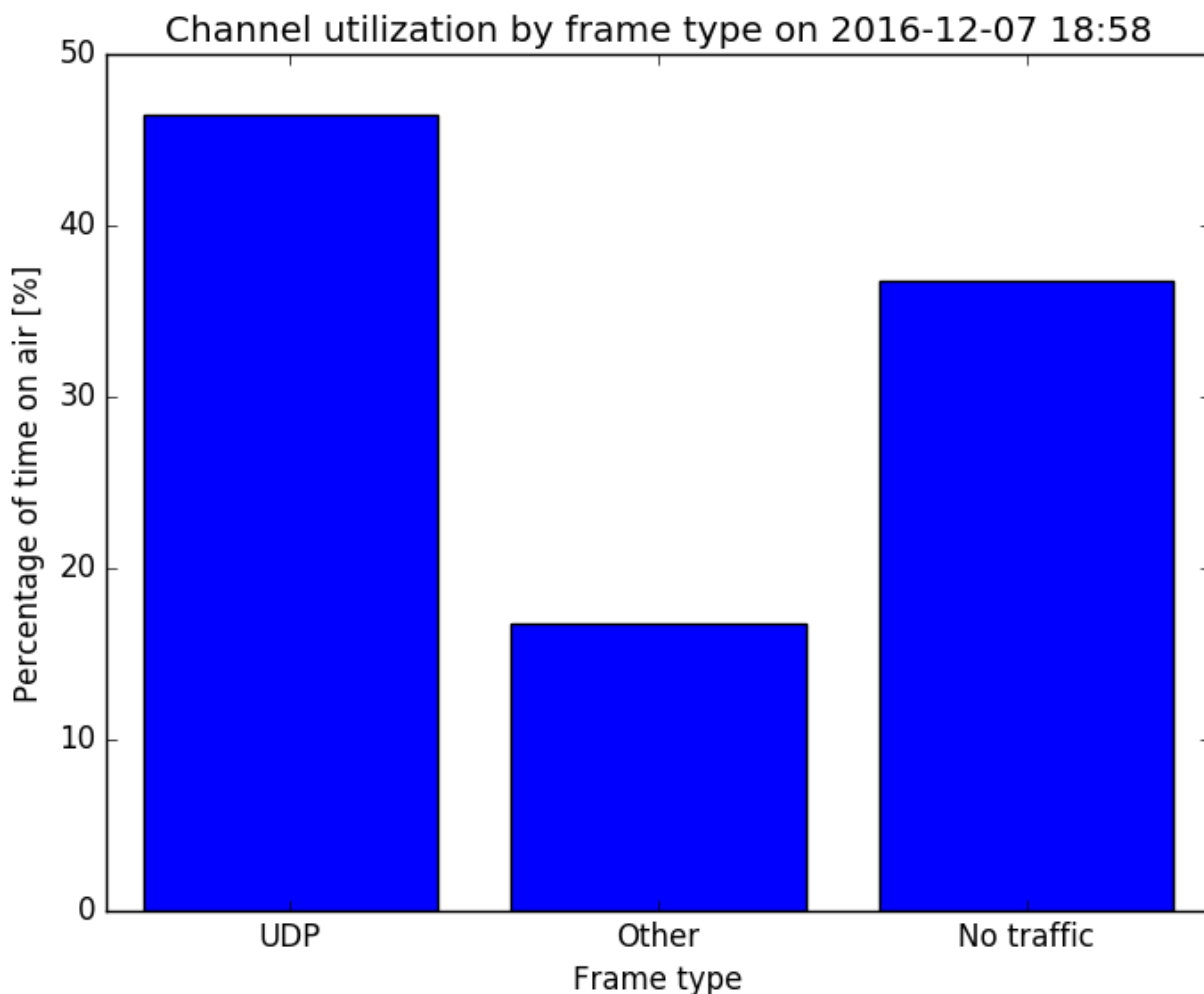
As for why the Channel utilization data we got from BSS Load elements (from the same trace as we computed the utilization manually) is about 6 seconds delayed we have no confident answer. We double checked the received tcpdump data in Wireshark and we made no mistake plotting the data. We can only assume that the Channel utilization data in BSS Load elements transmitted by other networks is late

because they average it, which causes the delay. Averaging seems a very likely answer when we look at the graph, as it has much lower spikes and variances and is a lot smoother.

## d) Medium utilization by frame type

---

### Plot



### Conclusion

Here as well, we excluded times for SIFS, DIFS, contention window and propagation delay and included ACK time for UDP packets. We filtered UDP packets with `data` in protocols, as we captured one UDP packet which was used for DHCP by another network, which was at the beginning of the trace and skewed the data quite considerably. Because we think that the purpose of this assignment is to get the distribution while an active UDP data transfer is in progress, we included the `data` filter for UDP packets to exclude the other packets.

As expected we see the majority of time taken up by UDP packets. Interestingly we can see quite a lot of time with no traffic. Some of that can be attributed to the omitted times in our frame time calculation, but we

don't expect that to be more than 10% or 15% in our case. That falls in line with the above graph for Medium utilization in time, where we can see that even though the iperf was throttled (actual bitrate reported was 7.5 Mbit/s) the channel was not used 100%, but instead has at least 20% non-busy time left. We expect this is by design.